

Overview

The Beta Distribution

Density

Examples

The Beta Distribution

Matthew Stephens

2017-01-25

Last updated: 2017-03-06

Code version: c7339fc

Overview

The purpose of this vignette is to introduce the Beta distribution. You should be familiar with basic concepts related to distributions before - e.g. maybe you have come across the normal distribution and a uniform distribution before, and understand what it would mean to talk about their mean, variance and density.

If you want more details you could look at Wikipedia (https://en.wikipedia.org/wiki/Beta_distribution).

The Beta Distribution

The Beta distribution is a distribution on the interval $[0, 1]$. Probably you have come across the $U[0, 1]$ distribution before: the uniform distribution on $[0, 1]$. You can think of the Beta distribution as a generalization of this that allows for some simple non-uniform distributions for values between 0 and 1.

The Beta distribution has two parameters, which we will call a and b . These two parameters determine the shape of the Beta distributions (just as the mean and variance determine the shape of the normal distribution).

Following the usual convention, we will write $X \sim Be(a, b)$ as shorthand for “ X has a Beta distribution with parameters a and b ”.

Density

If $X \sim Be(a, b)$ then the density of X is:

$$f_X(x) = \frac{1}{B(a, b)} x^{a-1} (1-x)^{b-1} \quad (x \in [0, 1]).$$

For those of you that are interested, $B(a, b)$ is known as the “beta function” and is given by the integral

$$B(a, b) = \int_0^1 x^{a-1} (1-x)^{b-1} dx.$$


This is where the beta distribution gets its name: its density involves the beta function. However, for this introduction you do not have to worry very much about what $B(a, b)$ is: think of it as a constant (in that it does not depend on x), that is included so that the density integrates to 1, as all densities must.

Because the Beta distribution is widely used, R has the built in function `dbeta` to compute this density. We will use this to look at some examples of the Beta distribution below.

Examples


First we will look at some examples for $a = b$, with both ≥ 1 :

```
p = seq(0,1, length=100)
plot(p, dbeta(p, 100, 100), ylab="density", type="l", col=4)
lines(p, dbeta(p, 10, 10), type="l", col=3)
lines(p, dbeta(p, 2, 2), col=2)
lines(p, dbeta(p, 1, 1), col=1)
legend(0.7,8, c("Be(100,100)", "Be(10,10)", "Be(2,2)", "Be(1,1)"),lty=c(
1,1,1,1),col=c(4,3,2,1))
```



Now non-equal values of a and b with both ≥ 1 :

```
p = seq(0,1, length=100)
plot(p, dbeta(p, 900, 100), ylab="density", type="l", col=4)
lines(p, dbeta(p, 90, 10), type="l", col=3)
lines(p, dbeta(p, 30, 70), col=2)
lines(p, dbeta(p, 3, 7), col=1)
legend(0.2,30, c("Be(900,100)", "Be(90,10)", "Be(30,70)", "Be(3,7)"), lty
=c(1,1,1,1), col=c(4,3,2,1))
```



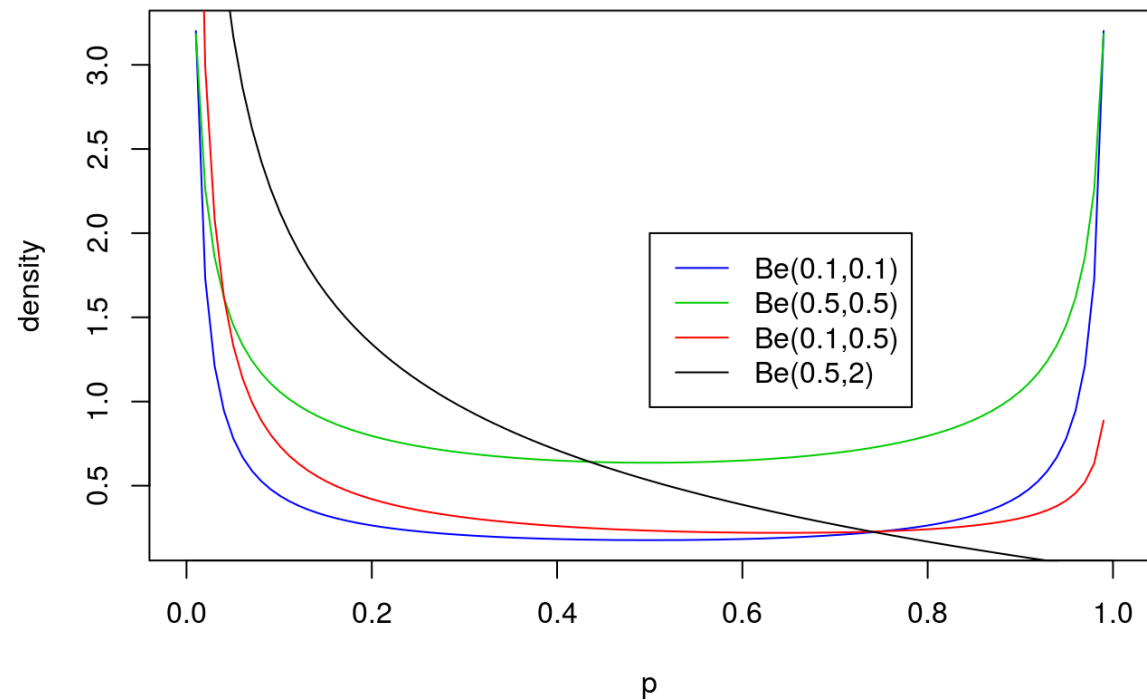
From these examples you should note the following:

- The distribution is roughly centered on $a/(a+b)$. Actually, it turns out that the mean is exactly $a/(a+b)$. Thus the mean of the distribution is determined by the *relative* values of a and b .
- The larger the values of a and b , the smaller the variance of the distribution about the mean.
- For moderately large values of a and b the distribution looks visually “kind of normal”, although unlike the normal distribution the Beta distribution is restricted to $[0,1]$.
- The special case $a = b = 1$ is the uniform distribution.

Values of $a, b < 1$

The parameters a and b can also be less than 1, but the distribution in this case starts to have a different kind of shape. Specifically if $a < 1$ then there is a peak at 0, and if $b < 1$ then there is a peak at 1 (so if both are < 1 then the distribution is U-shaped). Here are some examples:

```
p = seq(0,1, length=100)
plot(p, dbeta(p, 0.1, 0.1), ylab="density", type="l", col=4)
lines(p, dbeta(p, 0.5, 0.5), type="l", col=3)
lines(p, dbeta(p, 0.1, 0.5), col=2)
lines(p, dbeta(p, 0.5, 2), col=1)
legend(0.5,2, c("Be(0.1,0.1)","Be(0.5,0.5)","Be(0.1,0.5)", "Be(0.5,2)"
),lty=c(1,1,1,1),col=c(4,3,2,1))
```



Exercise

- Sketch what you think the $Be(5,5)$ and $Be(0.5,5)$ and $Be(500,200)$ distributions would look like. Check your sketches against the truth computed using `dbeta`.

Session information

```
sessionInfo()
```

```
R version 3.3.2 (2016-10-31)
Platform: x86_64-pc-linux-gnu (64-bit)
Running under: Ubuntu 14.04.5 LTS
```

```
locale:
```

```
[1] LC_CTYPE=en_US.UTF-8      LC_NUMERIC=C
[3] LC_TIME=en_US.UTF-8      LC_COLLATE=en_US.UTF-8
[5] LC_MONETARY=en_US.UTF-8  LC_MESSAGES=en_US.UTF-8
[7] LC_PAPER=en_US.UTF-8     LC_NAME=C
[9] LC_ADDRESS=C             LC_TELEPHONE=C
[11] LC_MEASUREMENT=en_US.UTF-8 LC_IDENTIFICATION=C
```

```
attached base packages:
```

```
[1] stats      graphics  grDevices  utils      datasets  methods    base
```

```
other attached packages:
```

```
[1] workflowr_0.4.0    rmarkdown_1.3.9004
```

```
loaded via a namespace (and not attached):
```

```
[1] backports_1.0.5 magrittr_1.5    rprojroot_1.2  htmltools_0.3.5
[5] tools_3.3.2     yaml_2.1.14    Rcpp_0.12.9    stringi_1.1.2
[9] knitr_1.15.1    git2r_0.18.0   stringr_1.2.0  digest_0.6.12
[13] evaluate_0.10
```

This site was created with R Markdown (<http://rmarkdown.rstudio.com>)

